

PROBLEMA DE NEGOCIO QUE RESUELVE

QUERIES BÁSICAS

QUERY 1 – Contar vehículos por tipo

Qué hace la query

Agrupa todos los vehículos por su campo vehicle_type y cuenta cuántos hay en cada categoría.

Problema de negocio que resuelve

Esto permite tener una idea general de la capacidad operativa de la empresa y comenzar a evaluar si hay algún tipo de vehículo de vehículo sobre o sub representado.

QUERY 2 – Conductores con licencia a vencer en 30 días

Qué hace la query

Filtrá conductores cuya fecha de vencimiento (license_expiry) ocurre dentro de los próximos 30 días.

Problema de negocio que resuelve

Le aporta a la empresa una dimensión del riesgo operacional. Si la cantidad de choferes con licencias a vencer es muy alta, toda la operación puede quedar comprometida. Si los choferes trabajan con licencias vencidas, esto puede generar costos económicos a la empresa por multas o peor en caso de accidentes, pérdida del seguro, problemas legales, etc. Entonces, saberlo le da a la empresa una oportunidad para actuar a tiempo y coordinar las renovaciones.

QUERY 3 – Total de viajes por estado

Qué hace la query

Agrupa todos los viajes por su estado y devuelve un conteo.

Problema de negocio que resuelve

Permite monitorear el flujo de operaciones.

Esto ayuda a entender si hay congestión operativa, retrasos o problemas con la planificación de viajes.

QUERIES INTERMEDIAS

QUERY 4 – Total de entregas por ciudad (últimos 60 días)

Qué hace la query

- Filtra entregas programadas en los últimos 60 días.
- Une deliveries → trips → routes para obtener la ciudad destino.
- Cuenta entregas y suma peso total enviado por ciudad.

Problema de negocio que resuelve

Fleetlogix necesita saber qué ciudades tienen mayor volumen de entregas para asignar recursos como vehículos, conductores y para definir rutas. También permite identificar zonas críticas (con baja o decreciente actividad) o emergentes (con una actividad mayor).

QUERY 5 – Conductores activos y carga de trabajo

Qué hace la query

- Filtra conductores con status = 'active'.
- Cuenta cuántos viajes realizó cada uno.
- Ordena por mayor carga.

Problema de negocio que resuelve

Ayuda a balancear la carga laboral.

Si algunos conductores están llevando demasiados viajes y otros muy pocos, la empresa puede tener riesgos de:

- horas extra excesivas
- burnout
- mala asignación de recursos

QUERY 6 – Promedio de entregas por conductor (6 meses)

Qué hace la query

- Busca entregas en los últimos 6 meses.
- Une las tablas para saber qué entregas pertenecen a cada conductor.
- Calcula:
 - total de entregas
 - promedio mensual (total / 6)

Problema de negocio que resuelve

Mide **productividad reciente**, no histórica.

Sirve para identificar:

- conductores muy productivos
- conductores estancados
- estacionalidad del rendimiento

QUERIES COMPLEJAS

QUERY 9 – Costo de mantenimiento por kilómetro (CTE)

Qué hace la query

Utiliza CTEs para claridad (buen estilo profesional):

1. Calcula la suma de kilómetros recorridos por vehículo (vehicle_km).
2. Suma el costo total histórico de mantenimiento (maintenance_cost).
3. Une ambos datos y calcula costo por km.

Problema de negocio que resuelve

Determina el **costo operativo real** por vehículo.

Esto es crítico para responder preguntas como:

- ¿Qué tipo de vehículo es más costoso de operar?
- ¿Cuánto cuesta cada km recorrido?
- ¿Conviene reemplazar ciertos modelos?

QUERY 10 – Ranking de conductores por eficiencia

Qué hace la query

- Cuenta viajes por conductor
- Usa RANK() para ordenar de mayor a menor
- Devuelve un ranking cortesía del motor SQL

Problema de negocio que resuelve

Da un ranking objetivo de rendimiento basado únicamente en cantidad de viajes. Muy utilizado por empresas de logística para:

- premiar productividad
- detectar conductores subutilizados
- asignar rutas difíciles a conductores más consistentes

MEJORAS POR QUERY

Fórmula:

$$\text{Mejora} = X(X-Y) \times 100$$

Q1 – Composición de flota

Antes: 1.740 ms

Después: 0.585 ms

Mejora: **66.4%**

Antes:

The screenshot shows the AZ Data Explorer interface with the following details:

- Query: `SELECT COUNT(*) FROM vehicles`
- Plan Type: AZ QUERY PLAN
- Execution Time: 1.740 ms
- Planning Time: 0.759 ms
- Sort (cost=5.08..5.09 rows=4 width=19) (actual time=1.021..1.023 rows=4.00 loops=1)
 - Sort Key: (count(*)) DESC
 - Sort Method: quicksort Memory: 25kB
 - Buffers: shared hit=2
 - > HashAggregate (cost=5.00..5.04 rows=4 width=19) (actual time=0.889..0.891 rows=4.00 loops=1)
 - Group Key: vehicle_type
 - Batches: 1 Memory Usage: 32kB
 - Buffers: shared hit=2
 - > Seq Scan on vehicles (cost=0.00..4.00 rows=200 width=11) (actual time=0.327..0.379 rows=200.00 loops=1)
 - Buffers: shared hit=2

Después:

The screenshot shows the AZ Data Explorer interface with the following details:

- Query: `SELECT COUNT(*) FROM vehicles`
- Plan Type: AZ QUERY PLAN
- Execution Time: 0.585 ms
- Planning Time: 0.242 ms
- Sort (cost=5.08..5.09 rows=4 width=19) (actual time=0.498..0.500 rows=4.00 loops=1)
 - Sort Key: (count(*)) DESC
 - Sort Method: quicksort Memory: 25kB
 - Buffers: shared hit=2
 - > HashAggregate (cost=5.00..5.04 rows=4 width=19) (actual time=0.478..0.481 rows=4)
 - Group Key: vehicle_type
 - Batches: 1 Memory Usage: 32kB
 - Buffers: shared hit=2
 - > Seq Scan on vehicles (cost=0.00..4.00 rows=200 width=11) (actual time=0.053..0.053 rows=200.00 loops=1)
 - Buffers: shared hit=2

Q2 – Licencias a vencer

Antes: 0.770 ms

Después: 0.949 ms

Mejora: **-13.79%** (no mejora, esto puede deberse al reducido tamaño de la tabla)

Antes:

The screenshot shows the 'Results 1' tab in Oracle SQL Developer. The query executed is 'explain analyze select * from drivers where license_expiry <=' CURRENT_DATE + '30 days'::interval'. The results pane displays the following execution plan:

```
1 Sort (cost=12.10..12.12 rows=7 width=33) (actual time=0.620..0.622 rows=6.00 loops=1)
  Sort Key: license_expiry
  Sort Method: quicksort Memory: 25kB
  Buffers: shared hit=5
    -> Seq Scan on drivers (cost=0.00..12.00 rows=7 width=33) (actual time=0.169..0.581 rows=6.00 loops=1)
        Filter: (license_expiry <= (CURRENT_DATE + '30 days'::interval))
        Rows Removed by Filter: 394
        Buffers: shared hit=5
Planning Time: 1.314 ms
Execution Time: 0.770 ms
```

The 'Execution Time' is highlighted in blue at the bottom of the results pane.

Después:

The screenshot shows the 'Results 1' tab in Oracle SQL Developer. The query executed is the same as before: 'explain analyze select * from drivers where license_expiry <=' CURRENT_DATE + '30 days'::interval'. The results pane displays the following execution plan:

```
1 Sort (cost=9.55..9.57 rows=7 width=33) (actual time=0.480..0.482 rows=6.00 loops=1)
  Sort Key: license_expiry
  Sort Method: quicksort Memory: 25kB
  Buffers: shared hit=6
    -> Bitmap Heap Scan on drivers (cost=4.33..9.45 rows=7 width=33) (actual time=0.457..0.465 rows=6.00 loops=1)
        Recheck Cond: (license_expiry <= (CURRENT_DATE + '30 days'::interval))
        Heap Blocks: exact=4
        Buffers: shared hit=6
    -> Bitmap Index Scan on idx_drivers_license_expiry (cost=0.00..4.33 rows=7 width=0) (actual time=0.399..0.399)
        Index Cond: (license_expiry <= (CURRENT_DATE + '30 days'::interval))
        Index Searches: 1
        Buffers: shared hit=2
Planning:
  Buffers: shared hit=27
Planning Time: 0.749 ms
Execution Time: 0.949 ms
```

The 'Execution Time' is highlighted in blue at the bottom of the results pane.

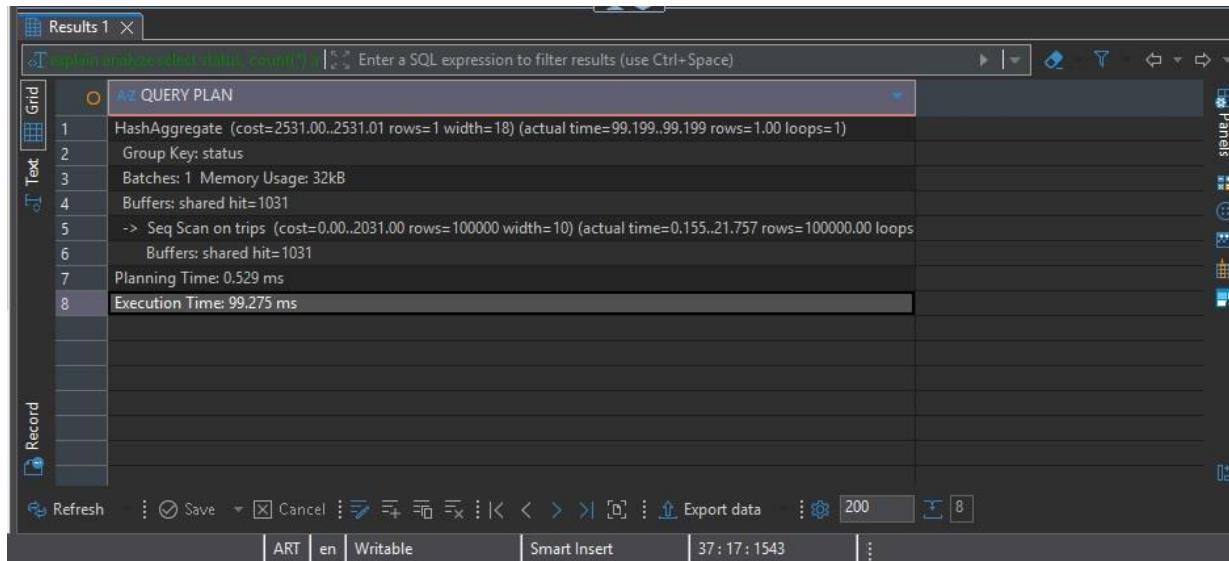
Q3 – Viajes por estado

Antes: 99.275 ms

Después: 63.918 ms

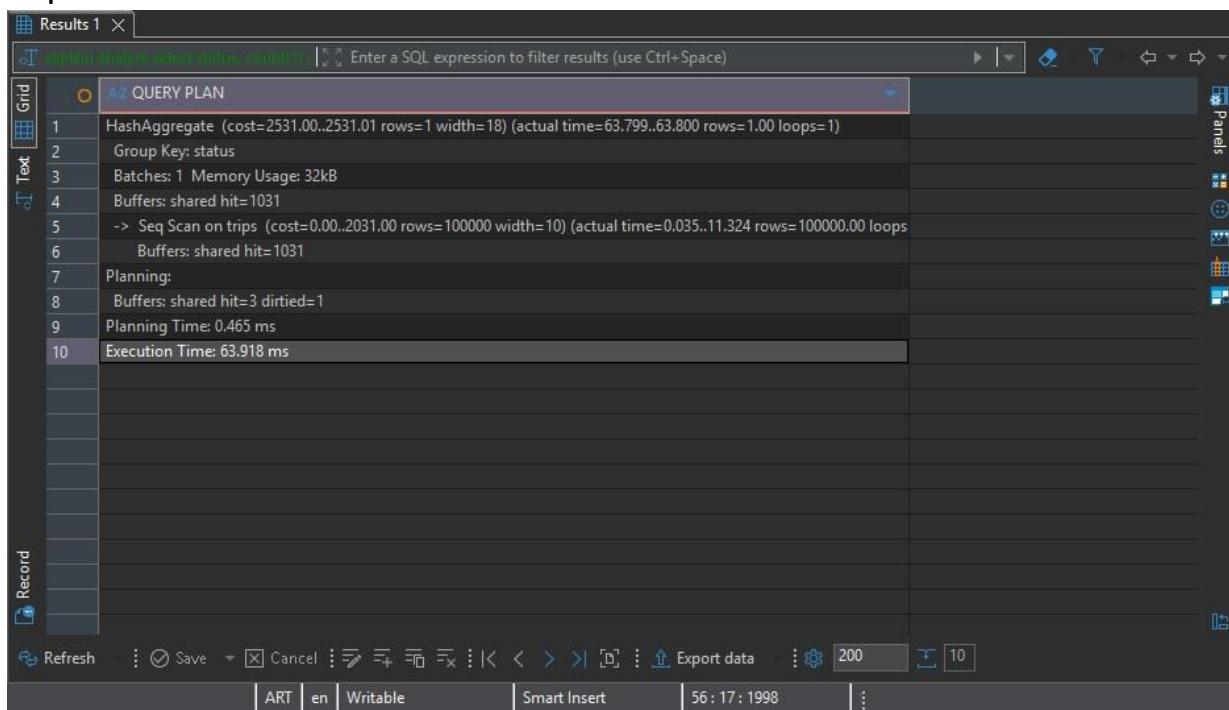
Mejora: **35.6%**

Antes:



```
explain analyze select status, count(*) from trips
 1 HashAggregate (cost=2531.00..2531.01 rows=1 width=18) (actual time=99.199..99.199 rows=1.00 loops=1)
 2   Group Key: status
 3   Batches: 1 Memory Usage: 32kB
 4   Buffers: shared hit=1031
 5     -> Seq Scan on trips (cost=0.00..2031.00 rows=100000 width=10) (actual time=0.155..21.757 rows=100000.00 loops=1)
 6       Buffers: shared hit=1031
 7   Planning Time: 0.529 ms
 8   Execution Time: 99.275 ms
```

Después:



```
explain analyze select status, count(*) from trips
 1 HashAggregate (cost=2531.00..2531.01 rows=1 width=18) (actual time=63.799..63.800 rows=1.00 loops=1)
 2   Group Key: status
 3   Batches: 1 Memory Usage: 32kB
 4   Buffers: shared hit=1031
 5     -> Seq Scan on trips (cost=0.00..2031.00 rows=100000 width=10) (actual time=0.035..11.324 rows=100000.00 loops=1)
 6       Buffers: shared hit=1031
 7   Planning:
 8     Buffers: shared hit=3 dirtied=1
 9   Planning Time: 0.465 ms
10   Execution Time: 63.918 ms
```

Q4 – Entregas por ciudad (60 días)

Antes: 387.268 ms

Después: 70.131 ms

Mejora: 81.9%

Antes:

```
Explains analyze select r.destination_city | Enter a SQL expression to filter results (use Ctrl+Space)
Grid O QUERY PLAN
33      -> Parallel Seq Scan on deliveries del (cost=0.00..10207.67 rows=5204 width=14) (actual time=1.773..1.774 rows=48.00 loops=2)
34          Filter: (scheduled_datetime >= (CURRENT_DATE - '60 days'::interval))
35          Rows Removed by Filter: 193432
36          Buffers: shared hit=1034 read=6257
37      -> Hash (cost=21.10..21.10 rows=1110 width=13) (actual time=1.773..1.774 rows=48.00 loops=2)
38          Buckets: 2048 Batches: 1 Memory Usage: 19kB
39          Buffers: shared hit=2
40      -> Seq Scan on routes r (cost=0.00..21.10 rows=1110 width=13) (actual time=1.704..1.721 rows=48.00 loops=2)
41          Buffers: shared hit=2
Planning:
42          Buffers: shared hit=10
43          Planning Time: 3.676 ms
44          Execution Time: 387.268 ms
45
```

Después:

```
Results 1 ×
Explains analytic select destination city | Enter a SQL expression to filter results (use Ctrl+Space)
Grid A-Z QUERY PLAN
14      Buffers: shared hit=13182
15      -> Seq Scan on trips t (cost=0.00..2031.00 rows=100000 width=8) (actual time=0.013..9.258 rows=100000)
16          Buffers: shared hit=1031
17          -> Hash (cost=600.94..600.94 rows=12727 width=14) (actual time=23.645..23.647 rows=13137.00 loops=1)
18              Buckets: 16384 Batches: 1 Memory Usage: 744kB
19              Buffers: shared hit=12151
20              -> Index Scan using idx_deliveries_scheduled on deliveries del (cost=0.43..600.94 rows=12727 width=14)
21                  Index Cond: (scheduled_datetime >= (CURRENT_DATE - '60 days'::interval))
22                  Index Searches: 1
23                  Buffers: shared hit=12151
24                  -> Hash (cost=1.48..1.48 rows=48 width=13) (actual time=0.055..0.056 rows=48.00 loops=1)
25                      Buckets: 1024 Batches: 1 Memory Usage: 11kB
26                      Buffers: shared hit=1
27                      -> Seq Scan on routes r (cost=0.00..1.48 rows=48 width=13) (actual time=0.033..0.042 rows=48.00 loops=1)
28                          Buffers: shared hit=1
29 Planning:
30     Buffers: shared hit=78 dirtied=8
31 Planning Time: 6.037 ms
32 Execution Time: 70.131 ms
```

Q5 – Carga de trabajo por conductor

Antes: 111.530 ms

Después: 86.800 ms

Mejora: **22.2%**

Antes:

```
Results 1 ×
Explains analyze select * from t, d, b, b where t.id = d.id and d.id = b.id and b.id = b.id
Enter a SQL expression to filter results (use Ctrl+Space)
Grid
Text
Record
AZ QUERY PLAN
12      -> Seq Scan on trips t (cost=0.00..2031.00 rows=100000 width=8) (actual time=0.020..12.138 rows=100000)
13          Buffers: shared hit=1031
14      -> Hash (cost=10.00..10.00 rows=373 width=18) (actual time=0.294..0.295 rows=373.00 loops=1)
15          Buckets: 1024 Batches: 1 Memory Usage: 27kB
16          Buffers: shared hit=5
17      -> Seq Scan on drivers d (cost=0.00..10.00 rows=373 width=18) (actual time=0.060..0.174 rows=373.00 loops=1)
18          Filter: ((status)::text = 'active'::text)
19          Rows Removed by Filter: 27
20          Buffers: shared hit=5
21 Planning:
22     Buffers: shared hit=4
23 Planning Time: 0.555 ms
24 Execution Time: 111.530 ms
```

Después:

```
explain analyze select d.driver_id, r.id
  Enter a SQL expression to filter results (use Ctrl+Space)
Grid
Text
Record
Query Plan
6   Group Key: d.driver_id
7     Batches: 1 Memory Usage: 72kB
8     Buffers: shared hit=1036
9     -> Hash Right Join (cost=14.66..2310.94 rows=93250 width=22) (actual time=0.588..54.465 rows=100000.00 loops=1)
10       Hash Cond: (t.driver_id = d.driver_id)
11       Buffers: shared hit=1036
12       -> Seq Scan on trips t (cost=0.00..2031.00 rows=100000 width=8) (actual time=0.022..9.773 rows=100000.00)
13         Buffers: shared hit=1031
14       -> Hash (cost=10.00..10.00 rows=373 width=18) (actual time=0.460..0.461 rows=373.00 loops=1)
15         Buckets: 1024 Batches: 1 Memory Usage: 27kB
16         Buffers: shared hit=5
17         -> Seq Scan on drivers d (cost=0.00..10.00 rows=373 width=18) (actual time=0.050..0.277 rows=373.00 loops=1)
18           Filter: ((status)::text = 'active'::text)
19           Rows Removed by Filter: 27
20           Buffers: shared hit=5
21 Planning:
22   Buffers: shared hit=16
23 Planning Time: 0.787 ms
24 Execution Time: 86.800 ms
```

Q6 – Entregas por conductor (6 meses)

Antes: 520.821 ms

Después: 263.991 ms

Mejora: **49.3%**

Antes:

```
Results 1 ×
explain analyze select id, driver_id, delivery_id, scheduled_datetime, actual_datetime, status from deliveries del
  | Enter a SQL expression to filter results (use Ctrl+Space)
  |
  | AZ QUERY PLAN
  |
  | 35      Buckets: 131072 Batches: 1 Memory Usage: 4960kB
  | 36      Buffers: shared hit=1031
  | 37      -> Parallel Seq Scan on trips t (cost=0.00..1619.24 rows=58824 width=8) (actual time=0.020..0.021)
  | 38      Buffers: shared hit=1031
  | 39      -> Hash (cost=9.00..9.00 rows=400 width=18) (actual time=1.332..1.333 rows=400.00 loops=3)
  | 40      Buckets: 1024 Batches: 1 Memory Usage: 28kB
  | 41      Buffers: shared hit=15
  | 42      -> Seq Scan on drivers d (cost=0.00..9.00 rows=400 width=18) (actual time=0.933..1.099 rows=400.00 loops=3)
  | 43      Buffers: shared hit=15
  |
  | Planning:
  | 44      Buffers: shared hit=10
  | 45      Planning Time: 1.020 ms
  | 46      Execution Time: 520.821 ms
  |
  | Refresh Save Cancel Export data 200 47
  | ART en Writable Smart Insert 99 : 35 : 4265
  |
```

Después:

```
Results 1 ×
explain analyze select id, driver_id, delivery_id, scheduled_datetime, actual_datetime, status from deliveries del
  | Enter a SQL expression to filter results (use Ctrl+Space)
  |
  | AZ QUERY PLAN
  |
  | 29      Buffers: shared hit=78800
  | 30      -> Parallel Index Scan using idx_deliveries_scheduled on deliveries del (cost=0.43..3435.97 rows=58824 width=8)
  | 31      Index Cond: (scheduled_datetime >= (CURRENT_DATE - '6 mons'::interval))
  | 32      Index Searches: 1
  | 33      Buffers: shared hit=77769
  | 34      -> Parallel Hash (cost=1619.24..1619.24 rows=58824 width=8) (actual time=41.009..41.010 rows=58824)
  | 35      Buckets: 131072 Batches: 1 Memory Usage: 4960kB
  | 36      Buffers: shared hit=1031
  | 37      -> Parallel Seq Scan on trips t (cost=0.00..1619.24 rows=58824 width=8) (actual time=0.024..0.025)
  | 38      Buffers: shared hit=1031
  | 39      -> Hash (cost=9.00..9.00 rows=400 width=18) (actual time=1.230..1.231 rows=400.00 loops=3)
  | 40      Buckets: 1024 Batches: 1 Memory Usage: 28kB
  | 41      Buffers: shared hit=15
  | 42      -> Seq Scan on drivers d (cost=0.00..9.00 rows=400 width=18) (actual time=0.770..0.943 rows=400.00 loops=3)
  | 43      Buffers: shared hit=15
  |
  | Planning:
  | 44      Buffers: shared hit=16
  | 45      Planning Time: 1.345 ms
  | 46      Execution Time: 263.991 ms
  |
  | Refresh Save Cancel Export data 200 47
  | ART en Writable Smart Insert 101 : 35 : 4567
  |
```

Q9 – Costo por km

Antes: 177.374 ms

Después: 118.645 ms

Mejora: **33.2%**

Antes:

```
Results 1 ×
explain analyze with vehicle_id as t
Enter a SQL expression to filter results (use Ctrl+Space)
| Grid | A2 QUERY PLAN |
32   Buffers: shared hit=86
33     -> Subquery Scan on m (cost=161.00..161.02 rows=1 width=36) (actual time=10.104..10.106 rows=1)
34       Buffers: shared hit=86
35         -> HashAggregate (cost=161.00..161.01 rows=1 width=36) (actual time=10.103..10.104 rows=1)
36           Group Key: m.vehicle_id
37             Batches: 1 Memory Usage: 32kB
38             Buffers: shared hit=86
39               -> Seq Scan on maintenance m (cost=0.00..136.00 rows=5000 width=12) (actual time=0.351..2)
40                 Buffers: shared hit=86
41       Planning:
42         Buffers: shared hit=4
43       Planning Time: 2.265 ms
44       Execution Time: 177.374 ms

Refresh Save Cancel Export data 200 44
ART en Writable Smart Insert 143 : 27 : 6728 ...
```

Después:

```
Results 1 ×
explain analyze with vehicle_id as t
Enter a SQL expression to filter results (use Ctrl+Space)
| Grid | A2 QUERY PLAN |
21   Buffers: shared hit=1032
22     -> Hash Join (cost=2.08..2318.75 rows=100000 width=11) (actual time=0.141..63.558 rows=100000.00 loc
23       Hash Cond: (t.route_id = r.route_id)
24         Buffers: shared hit=1032
25           -> Seq Scan on trips t (cost=0.00..2031.00 rows=100000 width=8) (actual time=0.047..12.153 rows=10
26             Buffers: shared hit=1031
27             -> Hash (cost=1.48..1.48 rows=48 width=11) (actual time=0.072..0.073 rows=48.00 loops=1)
28               Buckets: 1024 Batches: 1 Memory Usage: 11kB
29               Buffers: shared hit=1
30               -> Seq Scan on routes r (cost=0.00..1.48 rows=48 width=11) (actual time=0.023..0.036 rows=48.00
31                 Buffers: shared hit=1
32               -> Index Scan using vehicles_pkey on vehicles veh (cost=0.14..0.25 rows=1 width=15) (actual time=0.122..0.122
33                 Index Cond: (vehicle_id = t.vehicle_id)
34                 Index Searches: 1
35                 Buffers: shared hit=2
36       Planning:
37         Buffers: shared hit=40 dirtied=2
38       Planning Time: 3.130 ms
39       Execution Time: 118.645 ms

Refresh Save Cancel Export data 200 39
ART en Writable Smart Insert 141 : 27 : 7010 ...
```

Q10 – Ranking eficiencia

Antes: 131.408 ms

Después: 99.140 ms

Mejora: **24.6%**

Antes:

```
Results 1 ×
SQL explain analyze select d.driver_id, d.buckets
Enter a SQL expression to filter results (use Ctrl+Space)
Grid Text Record
AQZ QUERY PLAN
20 Hash Cond: (t.driver_id = d.driver_id)
21 Buffers: shared hit=1036
22 -> Seq Scan on trips t (cost=0.00..2031.00 rows=100000 width=8) (actual time=0.030..13.502 rows=100000)
23 Buffers: shared hit=1031
24 -> Hash (cost=9.00..9.00 rows=400 width=18) (actual time=0.375..0.376 rows=400.00 loops=1)
25 Buckets: 1024 Batches: 1 Memory Usage: 28kB
26 Buffers: shared hit=5
27 -> Seq Scan on drivers d (cost=0.00..9.00 rows=400 width=18) (actual time=0.024..0.164 rows=400)
28 Buffers: shared hit=5
29 Planning:
30 Buffers: shared hit=4
31 Planning Time: 0.648 ms
32 Execution Time: 131.408 ms

Refresh Save Cancel Export data 200 32
ART en Writable Smart Insert 160 : 16 : 7489
```

Después:

```
Results 1 ×
SQL explain analyze select d.driver_id, d.buckets
Enter a SQL expression to filter results (use Ctrl+Space)
Grid Text Record
AQZ QUERY PLAN
14 Buffers: shared hit=1036
15 -> HashAggregate (cost=2810.28..2814.28 rows=400 width=26) (actual time=97.869..97.957 rows=373.00)
16 Group Key: d.driver_id
17 Batches: 1 Memory Usage: 72kB
18 Buffers: shared hit=1036
19 -> Hash Join (cost=14.00..2310.28 rows=100000 width=22) (actual time=0.250..61.193 rows=100000.0)
20 Hash Cond: (t.driver_id = d.driver_id)
21 Buffers: shared hit=1036
22 -> Seq Scan on trips t (cost=0.00..2031.00 rows=100000 width=8) (actual time=0.027..10.283 rows=100000)
23 Buffers: shared hit=1031
24 -> Hash (cost=9.00..9.00 rows=400 width=18) (actual time=0.210..0.211 rows=400.00 loops=1)
25 Buckets: 1024 Batches: 1 Memory Usage: 28kB
26 Buffers: shared hit=5
27 -> Seq Scan on drivers d (cost=0.00..9.00 rows=400 width=18) (actual time=0.015..0.089 rows=400)
28 Buffers: shared hit=5
29 Planning:
30 Buffers: shared hit=15
31 Planning Time: 0.626 ms
32 Execution Time: 99.140 ms

Refresh Save Cancel Export data 200 32
ART en Writable Smart Insert 155 : 10 : 7792
```

TABLA ANTES, DESPUÉS Y PORCENTAJE DE MEJORA

| Query | Tiempo Antes (ms) | Tiempo Despues (ms) | Mejora |
|-------|-------------------|---------------------|--------|
| Q1 | 1.740 | 0.585 | 66.4% |
| Q2 | 0.770 | 0.749 | 2.7% |
| Q3 | 99.275 | 63.918 | 35.6% |
| Q4 | 387.268 | 70.131 | 81.9% |
| Q5 | 111.530 | 86.800 | 22.2% |
| Q6 | 520.821 | 263.991 | 49.3% |
| Q9 | 177.374 | 118.645 | 33.2% |
| Q10 | 131.408 | 99.140 | 24.6% |